

# A Framework to Transform In-Core GPU Algorithms to Out-of-Core Algorithms

## Supplemental material

Takahiro Harada\*  
Advanced Micro Devices, Inc.

### 1 Kernel Examples

Here we show a few kernel transformation examples. There are in core and out of core implementations. The modification made for the out of core implementation is highlighted in **bold**. These kernels are slightly simplified from the one used for the benchmark of the paper for illustration purpose.

#### 1.1 Kernel computing geometric normal and shading normal

##### 1. In core implementation

```
1 __kernel
2 void ComputeNormalKernel( __global Ray* gRays, __global Hit* gHits, __global HitNormal* gHitNormalOut,
3   __global Triangle* gTriStorage )
4 {
5     const int gIdx = GET_GLOBAL_IDX;
6     Triangle t;
7
8     if( hasHit( gHits[gIdx] ) )
9     {
10         t = gTriStorage[ gHits[gIdx].m_idx ];
11
12         const float4 ng = normalize3( cross3( t.v1-t.v0, t.v2-t.v0 ) );
13         const float4 hp = gRays[gIdx].getHitPoint();
14         const float4 bCrd = calcBaryCrd( hp, t.v0, t.v1, t.v2 );
15         const float4 ns = normalize3( bCrd.x * t.n0 + bCrd.y * t.n1 + bCrd.z * t.n2 );
16         gHitNormalOut[gIdx].m_ng = ng;
17         gHitNormalOut[gIdx].m_ns = ns;
18     }
19 }
```

##### 2. Out of core implementation

```
1 __kernel
2 void ComputeNormalKernel( __global Ray* gRays, __global Hit* gHits, __global HitNormal* gHitNormalOut,
3   VM.KERNEL_ARGS ) // ooc
4 {
5     const int gIdx = GET_GLOBAL_IDX;
6     Triangle t;
7
8     VMInitialize;
9
10    if( hasHit( gHits[gIdx] ) )
11    {
12        VMLoad( Triangle, gHits[gIdx].m_idx * sizeof(Triangle), &t, 0 );
13
14        const float4 ng = normalize3( cross3( t.v1-t.v0, t.v2-t.v0 ) );
15        const float4 hp = gRays[gIdx].getHitPoint();
16        const float4 bCrd = calcBaryCrd( hp, t.v0, t.v1, t.v2 );
17        const float4 ns = normalize3( bCrd.x * t.n0 + bCrd.y * t.n1 + bCrd.z * t.n2 );
18
19        gHitNormalOut[gIdx].m_ng = ng;
20        gHitNormalOut[gIdx].m_ns = ns;
21    }
22
23    VMFinalize;
24 }
```

---

\*e-mail:takahiro.harada@amd.com

## 1.2 Kernel casting rays to BVH to find the closest intersection

We used stackless BVH traversal to minimize the per-WI context. If stacked traversal is used, full stack needs to be loaded and stored.

### 1. In core implementation

```
1 __kernel
2 void RayCastKernel( __global Ray* gRays, __global Hit* gHits,
3   __global BvhNode* gBvh, __global Triangle* gTriStorage )
4 {
5     const int gIdx = GET_GLOBAL_IDX;
6
7     Triangle t;
8     float minDist = gHits[gIdx].m_fraction;
9     u32 minIdx = 0;
10    u32 nodeId = 0;
11    BvhNode node;
12    Ray ray = gRays[gIdx];
13
14    while( nodeId != BVH_TERMINATION )
15    {
16        node = gBvh[ nodeId ];
17
18        nodeId = NodeNext( &node );
19        Aabb aabb = NodeGetAabb( &node );
20        float ff = AabbIntersect( aabb, ray );
21        if( ff < minDist )
22        {
23            if( NodeIsLeaf( &node ) )
24            {
25                int faceIdx = NodeGetLeafData( &node );
26
27                t = gTriStorage[ faceIdx ];
28
29                float f = castRay( t.v0, t.v1, t.v2, ray );
30                if( f < minDist )
31                {
32                    minDist = f;
33                    minIdx = faceIdx;
34                }
35            }
36            else
37            {
38                nodeId = NodeGetChild0( &node );
39            }
40        }
41    }
42
43    if( minDist < gHits[gIdx].m_fraction )
44    {
45        gHits[gIdx].m_fraction = minDist;
46        gHits[gIdx].m_idx = minIdx;
47    }
48 }
```

### 2. Out of core implementation

```
1 __kernel
2 void RayCastKernel( __global Ray* gRays, __global Hit* gHits,
3   __global u64* gShapeOffsets, VM_KERNEL_ARGS ) // ooc
4 {
5     const int gIdx = GET_GLOBAL_IDX;
6
7     Triangle t;
8     float minDist = gHits[gIdx].m_fraction;
9     u32 minIdx = 0;
10    u32 nodeId = 0;
11    BvhNode node;
12    Ray ray = gRays[gIdx];
13
14    VMInitialize;
```

```

15
16 while( nodeIdX != BVH_TERMINATION )
17 {
18     VMLoad( BvhNode, gShapeOffsets[0] + nodeIdX * sizeof(BvhNode), &node, 0 );
19
20     nodeIdX = NodeNext( &node );
21     Aabb aabb = NodeGetAabb( &node );
22     float ff = AabbIntersect( aabb, ray );
23     if( ff < minDist )
24     {
25         if( NodeIsLeaf( &node ) )
26         {
27             int faceIdx = NodeGetLeafData( &node );
28
29             VMLoad( Triangle, gShapeOffsets[1] + faceIdx * sizeof(Triangle), &t, 1 );
30
31             float f = castRay( t.v0, t.v1, t.v2, ray );
32             if( f < minDist )
33             {
34                 minDist = f;
35                 minIdx = faceIdx;
36             }
37         }
38         else
39         {
40             nodeIdX = NodeGetChild0( &node );
41         }
42     }
43     pc = 0; // reset program counter
44 }
45
46 if( minDist < gHits[gIdx].m_fraction )
47 {
48     gHits[gIdx].m_fraction = minDist;
49     gHits[gIdx].m_idx = minIdx;
50 }
51
52 VMFinalize;
53 }

```

## 2 The Macro

Here we show the implementation of the macro used in the kernels shown above.

```

1 typedef struct
2 {
3     u64 m_baseAddr;
4     u64 m_dataSize;
5
6     u32 m_pageSize;
7     u32 m_nPages;
8     u32 m_padd;
9     u32 m_padd1;
10 } VMHeader;
11
12 typedef struct
13 {
14     u64 m_offset;
15     u32 m_isAvailable;
16     u32 m_time;
17 } VMPageTable;
18
19 u64 VMAAddr2Offset( u64 addr, VMHeader h ) { return (addr-h.m_baseAddr); }
20 u64 VMAAddr2PageIdx( u64 addr, VMHeader h ) { return VMAAddr2Offset( addr, h )/h.m_pageSize; }
21
22 #define VMLARGS __global VMPageTable* vmPt, __global char* vmStorage, const VMHeader vmHeader, const int
    vmFirst
23 #define VM_ARG_LIST vmPt, vmStorage, vmHeader, vmFirst
24 #define VM_KERNEL_ARGS __global u32* vmReqs, __global u32* vmNReqs, \

```

```

25  __global char* vmCtxt,\
26  VMARGS,\
27  const int vmNReqsMax
28
29  // Load the WI state and go where the WI was suspended
30  #define VMInitialize \
31  u64 vmAddr = 0;\
32  u32 pc = 0;\
33  if( !vmFirst )\
34  {\
35      LOAD_CTXT_MACRO;\
36      if( pc == VMPC_Finished ) return;\
37      switch( pc )\
38      {\
39          case 0: goto VMRESTART0;\
40          case 1: goto VMRESTART1;\
41          case 2: goto VMRESTART2;\
42          default: break;\
43      }\
44  }
45
46  // Save the WI state and append the page request if any
47  #define VMFinalize \
48  pc = VMPC_Finished;\
49  SAVE_CTXT_MACRO;\
50  return;\
51  SAVE_CTXT:\
52  {\
53      u32 o = AtomInc( vmNReqs[0] );\
54      AtomInc( vmReqs[(vmAddr-vmHeader.m_baseAddr)/vmHeader.m_pageSize] );\
55  }\
56  SAVE_CTXT_MACRO;
57
58  // Request memory at addr
59  #define VMRequest(TYPE, addr, dst, VMARGS) \
60  bool request##TYPE( u64 addr, TYPE* dst, VMARGS )\
61  {\
62      u64 o = VMAddr2Offset( addr, vmHeader );\
63      u64 pidx = VMAddr2PageIdx( addr, vmHeader );\
64      VMPageTable p = vmPt[pidx];\
65      if( !p.m_isAvailable )\
66          return false;\
67      u64 d = o - vmHeader.m_pageSize * pidx;\
68      (*dst) = *((__global TYPE*)&vmStorage[ p.m_offset + d ]);\
69      return true;\
70  }
71
72  VMRequest(float, addr, dst, VMARGS)
73
74  // Load data through the VM
75  #define VMLoad( TYPE, address, dst, LOADCOUNTER ) \
76  VMRESTART##LOADCOUNTER: \
77  {vmAddr = address;\
78  if( !request##TYPE( vmAddr, dst, VM_ARG_LIST ) ) \
79  { \
80      pc = LOADCOUNTER; \
81      goto SAVE_CTXT; \
82  }}

```